

Regression Analysis: Use the diabetes data set from UCI and Pima Indians Diabetes data set for performing the following:

a. Univariate analysis: Frequency, Mean, Median, Mode, Variance, Standard Deviation, Skewness and Kurtosis

b. Bivariate analysis: Linear and logistic regression modeling

c. Multiple Regression analysis

d. Also compare the results of the above analysis for the two data sets Dataset link:

<https://www.kaggle.com/datasets/uciml/pima-indians-diabetes-database>

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import r2_score, mean_squared_error, accuracy_score
import warnings
warnings.filterwarnings("ignore")
```

```
df=pd.read_csv(r"C:\Users\dell\Desktop\DMV and ML\ML Datasets\
diabetes.csv")
```

df

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

DiabetesPedigreeFunction Age Outcome

0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

df.describe()

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000
75%	6.000000	140.250000	80.000000	32.000000
max	17.000000	199.000000	122.000000	99.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
# univariate_analysis = pd.DataFrame({
#     'Frequency': df.count(),
#     'Mean': df.mean(),
#     'Median': df.median(),
#     'Mode': df.mode().iloc[0],
```

```
# 'Variance': df.var(),  
# 'Standard Deviation': df.std(),  
# 'Skewness': df.skew(),  
# 'Kurtosis': df.kurt()  
# }).T
```

```
# univariate_analysis
```

```
df.count()
```

Pregnancies	768
Glucose	768
BloodPressure	768
SkinThickness	768
Insulin	768
BMI	768
DiabetesPedigreeFunction	768
Age	768
Outcome	768

```
dtype: int64
```

```
df.skew()
```

Pregnancies	0.901674
Glucose	0.173754
BloodPressure	-1.843608
SkinThickness	0.109372
Insulin	2.272251
BMI	-0.428982
DiabetesPedigreeFunction	1.919911
Age	1.129597
Outcome	0.635017

```
dtype: float64
```

```
df.kurt()
```

Pregnancies	0.159220
Glucose	0.640780
BloodPressure	5.180157
SkinThickness	-0.520072
Insulin	7.214260
BMI	3.290443
DiabetesPedigreeFunction	5.594954
Age	0.643159
Outcome	-1.600930

```
dtype: float64
```

```
df.var()
```

```
Pregnancies      11.354056
Glucose          1022.248314
BloodPressure    374.647271
SkinThickness    254.473245
Insulin          13281.180078
BMI              62.159984
DiabetesPedigreeFunction  0.109779
Age              138.303046
Outcome          0.227483
dtype: float64
```

```
df.median()
```

```
Pregnancies      3.0000
Glucose          117.0000
BloodPressure     72.0000
SkinThickness     23.0000
Insulin           30.5000
BMI               32.0000
DiabetesPedigreeFunction  0.3725
Age               29.0000
Outcome           0.0000
dtype: float64
```

```
df.mode().iloc[0]
```

```
Pregnancies      1.000
Glucose          99.000
BloodPressure     70.000
SkinThickness     0.000
Insulin           0.000
BMI               32.000
DiabetesPedigreeFunction  0.254
Age               22.000
Outcome           0.000
Name: 0, dtype: float64
```

```
# x=df.iloc[:, :-1]
```

```
# y=df.iloc[:, -1]
```

```
x=df.drop("Outcome",axis=1)
```

```
y=df["Outcome"]
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=100)
```

```
model = LinearRegression()
model.fit(x_train,y_train)
y_pred_linear=model.predict(x_test)
```

```
print("Linear Regression R-squared:", r2_score(y_test, y_pred_linear))
print("MSE :", mean_squared_error(y_test, y_pred_linear))
```

```
Linear Regression R-squared: 0.19996694950228422
MSE : 0.1805775391851186
```

```
logistic = LogisticRegression()
logistic.fit(x_train, y_train)
y_pred_logistic = logistic.predict(x_test)
```

```
accuracy = accuracy_score(y_test, y_pred_logistic)
print(f"Logistic Regression Accuracy: {accuracy}")
```

```
Logistic Regression Accuracy: 0.7532467532467533
```

```
# X = df[['Age', 'BMI']]
# y = df['Outcome']
```

```
# model = LinearRegression()
# model.fit(X, y)
```

```
# # You can now analyze the coefficients, make predictions, and
evaluate the model's performance.
```

```
# new_data = pd.DataFrame({'Age': [40, 45], 'BMI': [30, 35]})
# predictions = model.predict(new_data)
# print("Predictions:", predictions)
```

```
# X = df.drop(["Outcome"], axis=1)
# y = df['Outcome']
```

```
# log_model = LogisticRegression()
# log_model.fit(X, y)
```

```
# new_data = pd.DataFrame({'Pregnancies': [5, 3], 'Glucose': [120,
160], 'BloodPressure': [70, 80], 'SkinThickness': [30, 20], 'Insulin':
[0, 40], 'BMI': [25.5, 28.6], 'DiabetesPedigreeFunction': [0.5, 0.4],
'Age': [35, 28]})
# probabilities = log_model.predict_proba(new_data)
# print("Probabilities of having diabetes:", probabilities)
```


standard deviation = calculate the dispersion of the dataset relative to its mean, calculated by

$$\text{Standard Deviation} = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n - 1}}$$

square root of variance

variance= Variance is a statistical measurement of the spread between numbers in a data set

standard deviation is the square root of the variance, and variance is the average of the squared difference of each data point from the mean.




skewness= Skewness is a measure of the asymmetry of a distribution (Skewness = Mean – Mode)

Positive Skewness (Right Skew)-- $\text{Mean} > \text{Median} > \text{Mode}$

Negative Skewness (Left Skew)-- $\text{Mean} < \text{Median} < \text{Mode}$

kurtosis: statistical measure used to describe the distribution of the observed data around mean



An R-Squared value shows how well the model predicts the outcome of the dependent variable. R-Squared values range from 0 to 1. An R-Squared value of 0 means that the model explains or predicts 0% of the relationship between the dependent and independent variables.

$$R^2 = 1 - \frac{RSS}{TSS}$$

R^2 = coefficient of determination

RSS = sum of squares of residuals

TSS = total sum of squares

$$\begin{aligned} r2_score &= 1 - \frac{\text{total_error_model}}{\text{total_error_baseline}} \\ &= 1 - \frac{\sum_{i=1}^N (\text{predicted}_i - \text{actual}_i)^2}{\sum_{i=1}^N (\text{average_value} - \text{actual}_i)^2} \end{aligned}$$

mean_squared_error -- average squared difference between the predicted values and the actual values in the dataset.(lower the better)

Mean

Error **Squared**

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

Univariate Analysis For each dataset, compute the following for each numeric column:

Frequency: Count occurrences of each unique value (particularly useful for categorical data).

Mean, Median, and Mode: Central tendency measures.

Variance and Standard Deviation: Spread measures to understand data variability.

Skewness: Measures asymmetry of the data distribution. Skewness > 0 indicates a right-skewed distribution; Skewness < 0 indicates a left-skewed distribution.

Kurtosis: Indicates the "tailedness" of the distribution. Kurtosis > 0 indicates a heavy-tailed distribution, while Kurtosis < 0 suggests a light-tailed distribution.

```
x=df[["Glucose"]]
y=df["Outcome"]

model3=LinearRegression()
model3.fit(x,y)
y_pred_model3=model3.predict(x)

mse = mean_squared_error(y, y_pred_model3)
r2 = r2_score(y, y_pred_model3)
print("Mean Squared Error:", mse)
print("R-squared:", r2)
```

```
Mean Squared Error: 0.17772834105264668
R-squared: 0.21769820124599804
```